

Computer Networks

Transport Layer

Presented by Hung Ba Ngo


Tháng 08/2015

Objectives


- Roles of the transport layer
- Services provided by the transport layer
- Connection establishment
- Connection release
- TCP and UDP protocols

Roles of Transport layer

- While Network layer provides a host-to-host communication, Transport layer provides a End point-to-End point communication
- End points are running applications
- Provides effective, reliable and cost savings packet transmission service for users

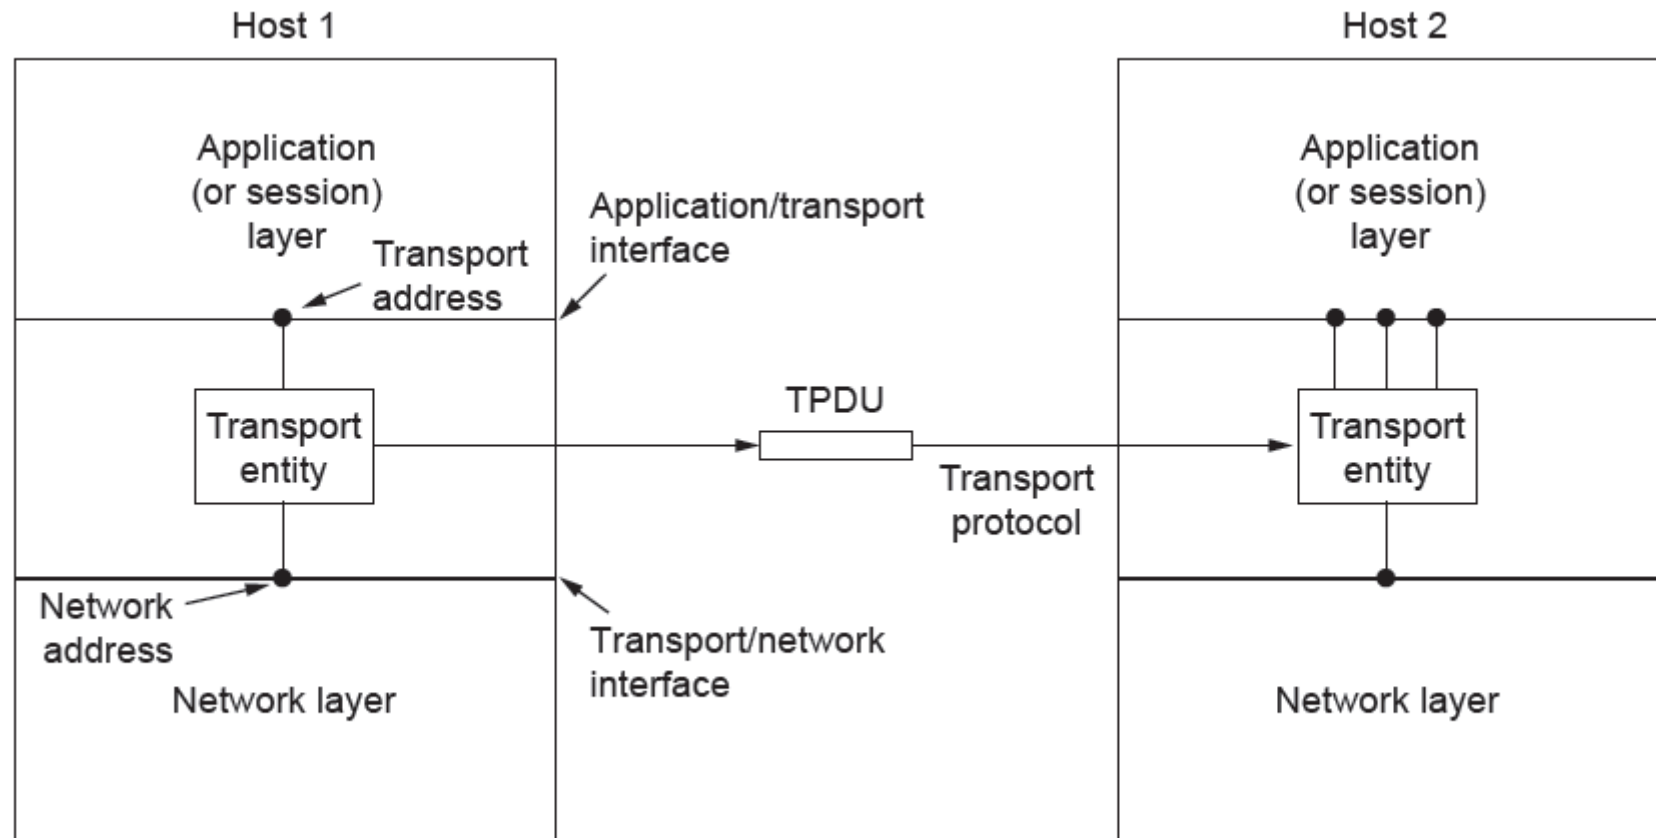


Environment of the ^(a) data link layer.



Environment of the ^(b) transport layer

Services Provided to the Upper Layers

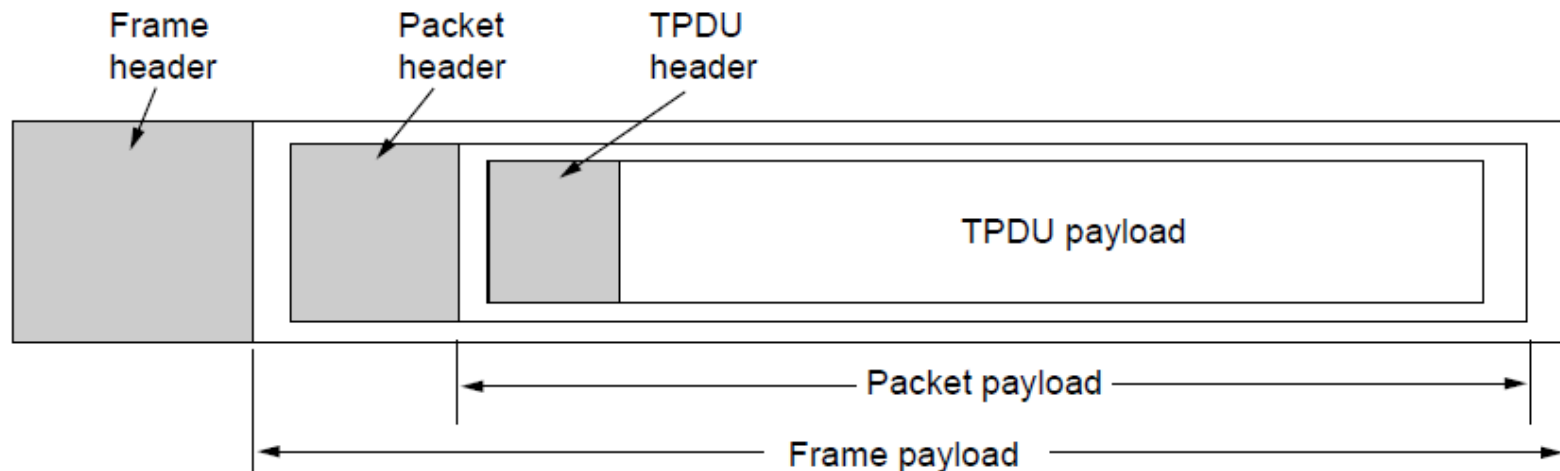


The network, transport, and application layers

Transport Service Primitives

| Primitive | Packet sent | Meaning |
|------------|--------------------|--|
| LISTEN | (none) | Block until some process tries to connect |
| CONNECT | CONNECTION REQ. | Actively attempt to establish a connection |
| SEND | DATA | Send information |
| RECEIVE | (none) | Block until a DATA packet arrives |
| DISCONNECT | DISCONNECTION REQ. | This side wants to release the connection |

The primitives for a simple transport service



Nesting of TPDUs, packets, and frames.

Berkeley Sockets

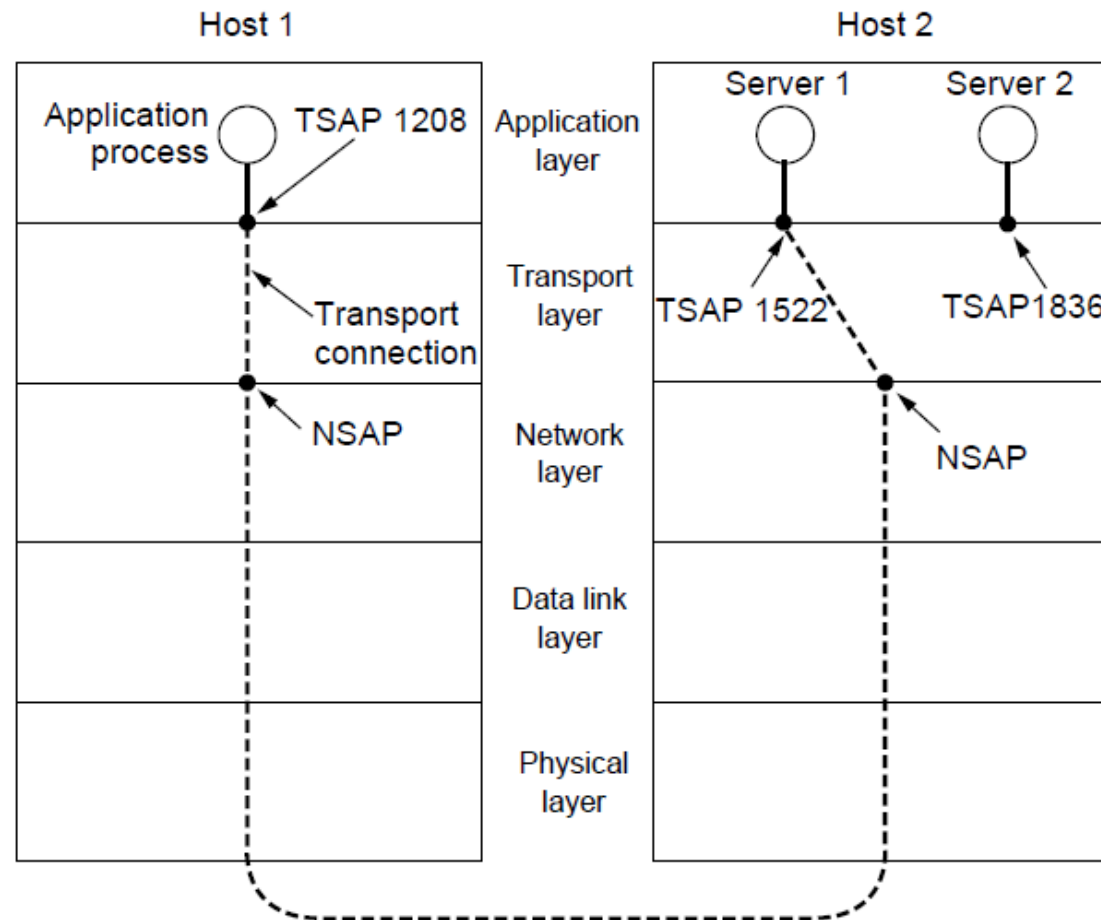
| Primitive | Meaning |
|-----------|---|
| SOCKET | Create a new communication end point |
| BIND | Associate a local address with a socket |
| LISTEN | Announce willingness to accept connections; give queue size |
| ACCEPT | Passively establish an incoming connection |
| CONNECT | Actively attempt to establish a connection |
| SEND | Send some data over the connection |
| RECEIVE | Receive some data from the connection |
| CLOSE | Release the connection |

The socket primitives for TCP


Elements of Transport Protocols

- Addressing
- Connection establishment
- Connection release
- Error control and flow control
- Multiplexing
- Crash recovery

Addressing




TSAPs, NSAPs, and transport connections




How a user process in host 1 establishes a connection with a mail server in host 2 via a process server.

Connection Establishment



Establishing a connection using a
three-way handshake.
CR denotes CONNECTION REQUEST.
Normal operation.

Connection Establishment

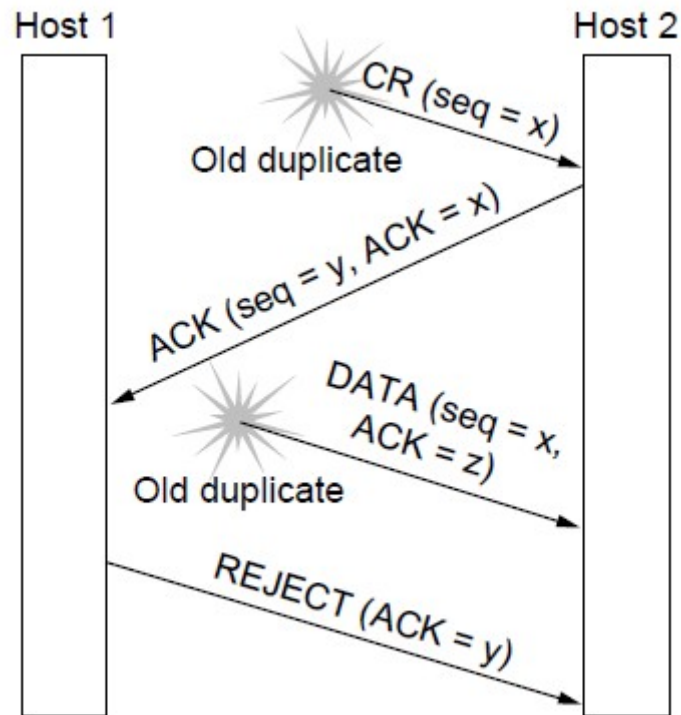


Establishing a connection using a
three-way handshake.

CR denotes CONNECTION REQUEST.

Old duplicate CONNECTION REQUEST appearing out of nowhere.

Connection Establishment




Establishing a connection using a
three-way handshake.

CR denotes CONNECTION REQUEST.


Duplicate CONNECTION REQUEST and duplicate ACK

Connection Release



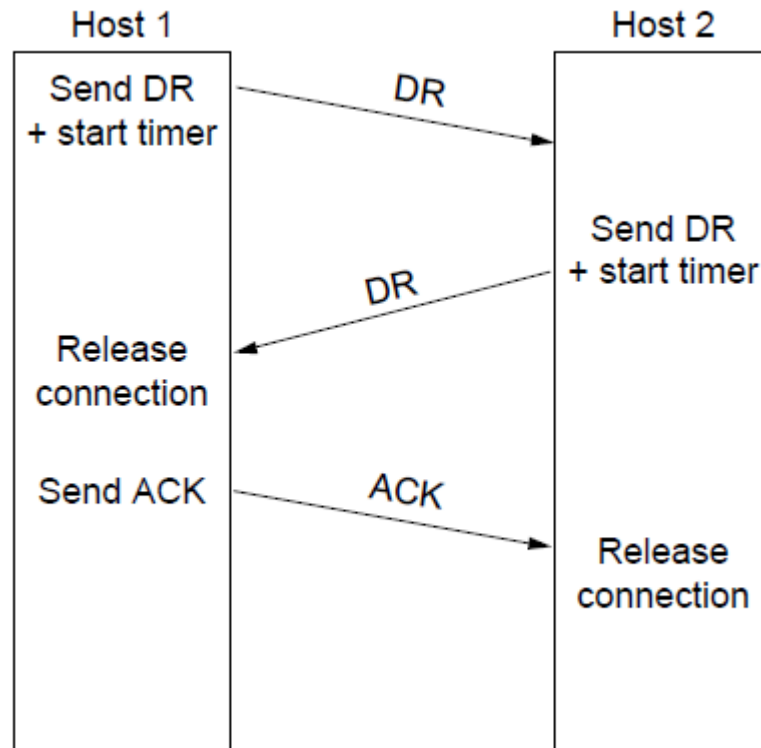
Abrupt disconnection with loss of data

Connection Release



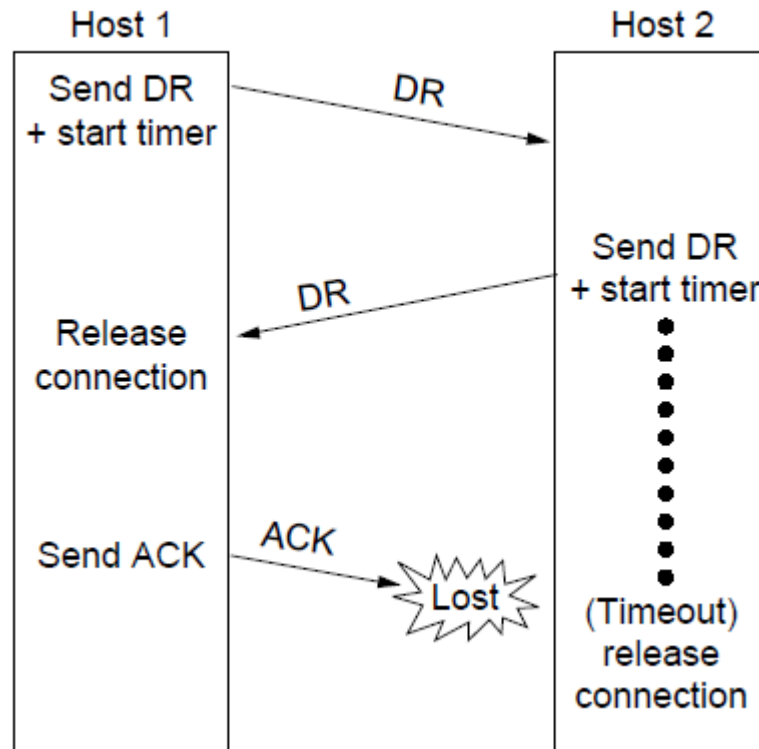
The two-army problem

Connection Release



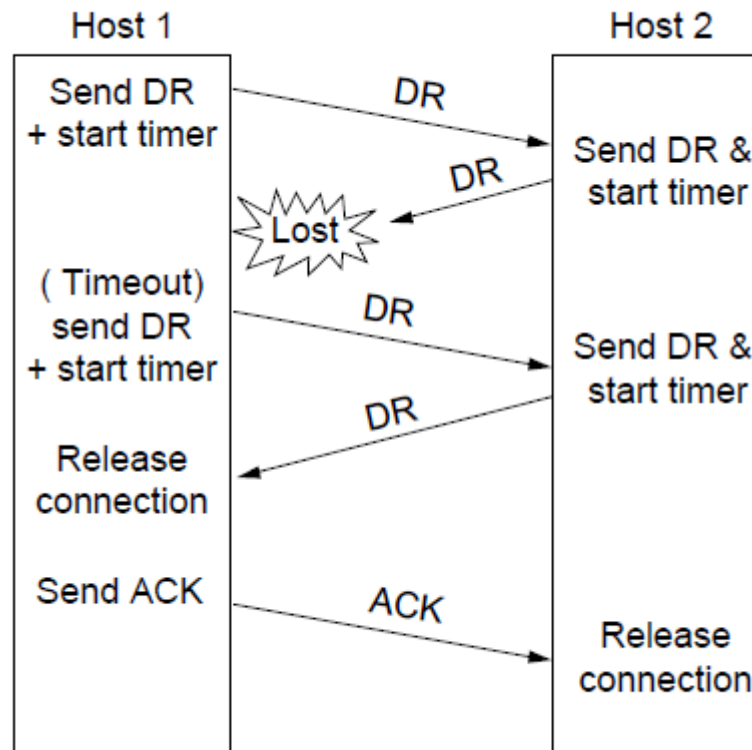
Four protocol scenarios for releasing a connection.
(a) Normal case of three-way handshake

Connection Release



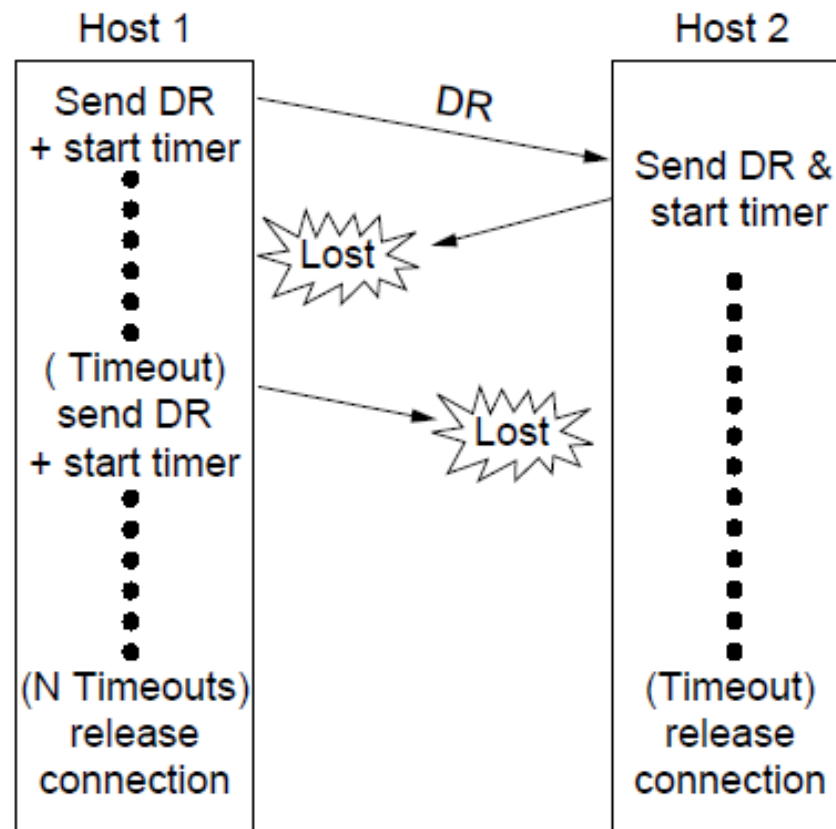
Four protocol scenarios for releasing a connection.
(b) Final ACK lost.

Connection Release



Four protocol scenarios for releasing a connection.
(c) Response lost

Connection Release



Four protocol scenarios for releasing a connection.
(d) Response lost and subsequent DRs lost.

Flow Control

- Using Sliding window protocol where the size of the sending window and the receiving window are different
- Need a scheme for allocating buffer dynamically

Flow Control

| | <u>A</u> | <u>Message</u> | <u>B</u> | <u>Comments</u> |
|----|----------|----------------------|----------|--|
| 1 | → | < request 8 buffers> | → | A wants 8 buffers |
| 2 | ← | <ack = 15, buf = 4> | ← | B grants messages 0-3 only |
| 3 | → | <seq = 0, data = m0> | → | A has 3 buffers left now |
| 4 | → | <seq = 1, data = m1> | → | A has 2 buffers left now |
| 5 | → | <seq = 2, data = m2> | ... | Message lost but A thinks it has 1 left |
| 6 | ← | <ack = 1, buf = 3> | ← | B acknowledges 0 and 1, permits 2-4 |
| 7 | → | <seq = 3, data = m3> | → | A has 1 buffer left |
| 8 | → | <seq = 4, data = m4> | → | A has 0 buffers left, and must stop |
| 9 | → | <seq = 2, data = m2> | → | A times out and retransmits |
| 10 | ← | <ack = 4, buf = 0> | ← | Everything acknowledged, but A still blocked |
| 11 | ← | <ack = 4, buf = 1> | ← | A may now send 5 |
| 12 | ← | <ack = 4, buf = 2> | ← | B found a new buffer somewhere |
| 13 | → | <seq = 5, data = m5> | → | A has 1 buffer left |
| 14 | → | <seq = 6, data = m6> | → | A is now blocked again |
| 15 | ← | <ack = 6, buf = 0> | ← | A is still blocked |
| 16 | ... | <ack = 6, buf = 4> | ← | Potential deadlock |

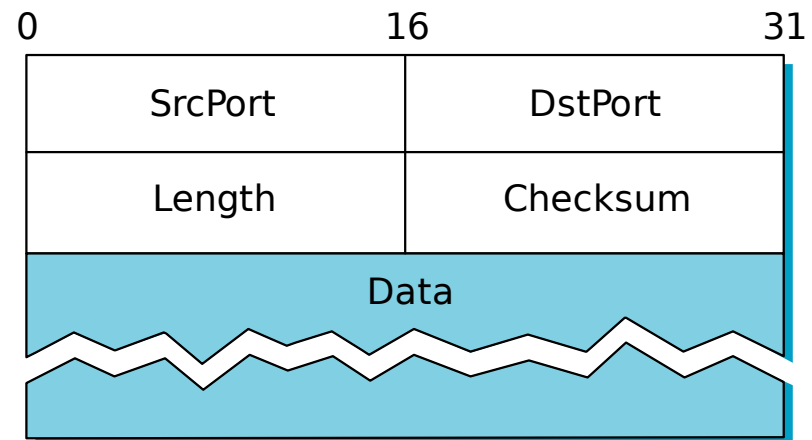
Dynamic buffer allocation. The arrows show the direction of transmission. An ellipsis (...) indicates a lost TPDU

Internet Protocol Suite

UDP & TCP

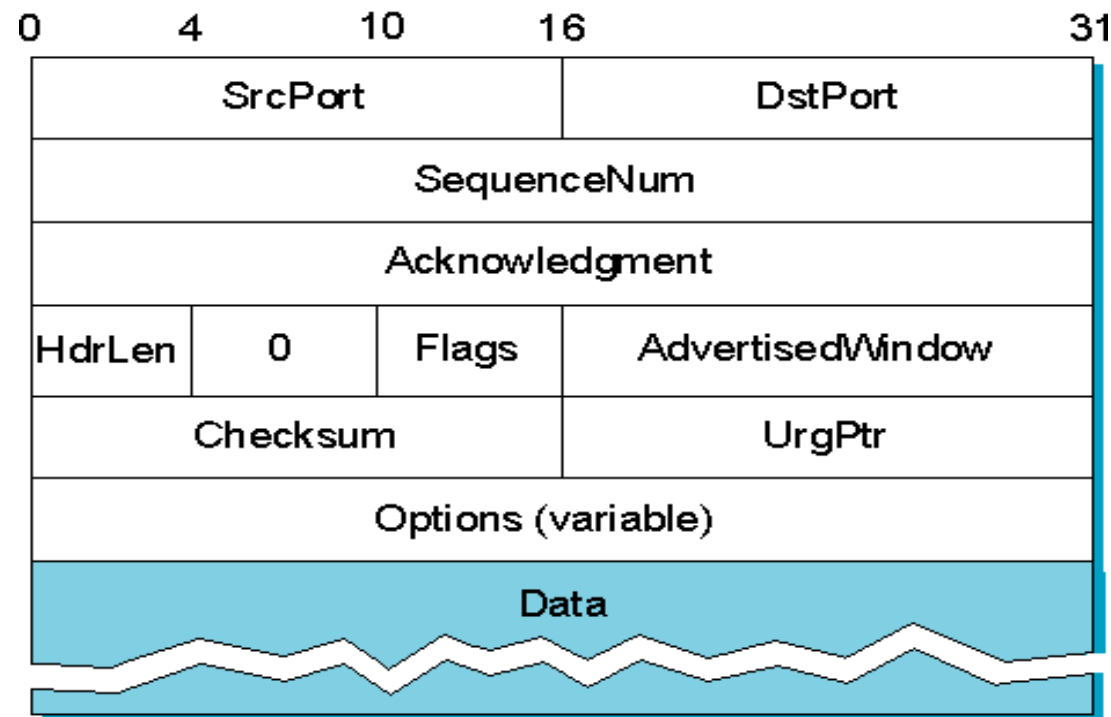
UDP - User Datagram Protocol

- A connectionless communication protocol
- Not required to establish a connection between two end points
- A UDP segment could appear at a destination at anytime
- A UDP segment contains enough information that can be used to forward the segment to the destination




TCP - Transmission Control Protocol

- A connection-oriented communication protocol
- Byte-oriented protocol



Flags = [SYN, FIN, RESET, PUSH, URG, ACK]

TCP – Connection establishment



- a) TCP connection establishment in the normal case
- b) Simultaneous connection establishment on both sides

Flow control with sliding window protocol

